

CS 111: Operating System Principles

Lab 1

Pipe Up_{3.0.0}

Can Aygun, Alexandre Tiard, Vishnu Vardhan Bachupally

Derivative document by: Jonathan Eyolfson

April 8, 2022

Due: April 22, 2022 @ 11:59 PM PT

In this lab you'll be writing the low level code performed by the pipe (|) operator in shells. Users will pass in executable names as command line arguments and you will execute each one in a new process (again similar to a shell). You're expected to understand your own implementation and test it yourself.

Additional APIs. Given you are just using the executable names, you may use the C library helper functions for `execve`, such as `execvp`. `execvp` will let you skip using string arrays (using C varargs), and it will also search for executables using the `PATH` environment variable. You may notice your program hanging waiting for input. You must call `close` on any file descriptors not explicitly used in your process. Failure to call `close` will inform the operating system you are not done with it, and it will never return end-of-file from a `read` system call.

Starting the lab. You'll receive all future labs through our repository (Lab 0 had you set it up as upstream). Run the following command to get the skeleton for Lab 1: `git pull upstream main`. You should be able to run `make` in the `lab1` directory to create a pipe executable, and then `make clean` to remove all binary files.

Files to modify. You should only be modifying `pipe.c` and `README.md` in the `lab1` directory.

Your task. You should execute the programs in `argv[1]`, ..., `argv[argc - 1]` as new processes. You also need to create a pipe between two subsequent processes. For example, a pipe should connect `argv[1]`'s standard output to `argv[2]`'s standard input (if there are at least two processes). The standard input of first new process must be the same as the standard input of the parent process (pipe). Also, the standard output of your last new process must be the same as the standard output of the parent process. You should be able to handle between 1 to at least 8 programs (more is okay). All standard errors should be the same as the parent's standard error. You do not need to handle passing additional command line arguments to every individual new process. All your processes should be a direct child of the original process that starts executing `main`. Finally, fill in your `README.md` so that you could use your program without having to use this document.

Errors. Your program should (of course) handle errors from all function calls you make. Your program should exit with the proper `errno` of the failing call. It is okay to exit as soon as you find an error, without any error recovery. If there are no programs as command line arguments, your program should exit with `errno EINVAL` (invalid argument). Your program should work with a single program as a command line argument. Your program should not create any orphan processes (you must wait).

Tips. You should come up with smaller subtasks yourself. One approach may be to execute one program from the command line, then multiple programs independently. Afterwards set up your pipe between two processes, then multiple processes. Start small then work big.

Example output. Your output should be the same as if you were to use | between each program in your shell.

```
> ./pipe ls
Makefile  pipe  pipe.c  pipe.o  README.md
> ./pipe ls cat wc
5         5         38
```

The last command should have the same output of: `ls | cat | wc`.

Testing. There are a set of basic test cases given to you (they'll be released a few days after the lab). We'll withhold more advanced tests which we'll use for grading. Part of programming is coming up with tests yourself. To run the provided test cases please run the following command in your lab directory:

```
python -m unittest
```

Submission. Simply push your code using `git push origin main` (or simply `git push`) to submit it. *You need to create your own commits to push, you can use as many as you'd like.* You'll need to use the `git add` and `git commit` commands. You may push as many commits as you want, your latest commit that modifies the lab files counts as your submission. For late days we will look at the timestamp on our server. We will never use your commit times (or file access times) as proof of submission, only when you push your code to the course Git server.

Please check <https://sefer.cs.ucla.edu/cs111/grades/> to see your status. You're expected to properly merge in upstream code without rebasing. Note that the website only updates your lab modification status if you've merged the latest code.